

Syllabus etc

Monday, August 22, 2016 10:02 AM

CS 3600

Mark Riedl riedl@gatech.edu

TSRB 228

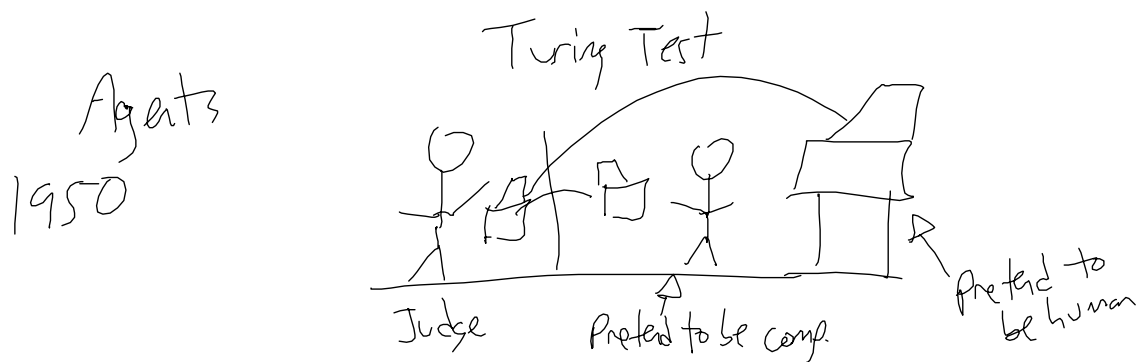
OH: Tues 2:30 - 4:30p

www.cc.gatech.edu/~riedl/classes/2016/cs3600/PREFREQSData structures - linked lists, trees, graphsComp Complexity - Big-OhNP-Hard $O(K^n)$ Programming Mojo - algorithms + systems
- PythonLogistics

- Lectures
 - Reading - schedule
 - HW assignments - 4 60%
 - HW
 - Exercises 0%
- Extra Credit
it
1.5%
max 2

- Tests - Midterm / Final 20% each

An entity is intelligent if it performs behaviors that a person might reasonably believe requires intelligence.



Strong AI:

- emulates humans
- broad & general
- beating the T.T.

Weak AI:

- Solving specific problems
- Narrow
- Automation
- Super-Intelligent

AI is a moving target

Machine Learning

- AI that self-improves
- Finding patterns in data & acting on it

Think like a human	Think rationally, ①
Act like a human	Act rationally

Think Rational.

Logic

Joe is a human
All humans are mortal
Mortal humans can be killed
Q: Can I kill Joe?

Formal representation

Alg. that manipulate facts & prove theorems

Q: Can I kill ~~Bob~~?
Elmo?

Chess checkers

Acting Rationally:

- Don't care how I get the answer
- Am I acting optimally
- Choose actions that lead to best outcomes
- Don't have right all the time

Memories

Think like a human:

- Cognitive science
- AI program is a hypothesis about the human mind

Acting like a human:


- Beating fu Turing Test
- Chatbots

History

Tuesday, August 23, 2016 6:26 AM



1200s
Automaton



~~1700s~~

1770

Computer ~~1780~~
1800

• **Early vision: Neanderthal AI (pre 1960s)**

- Model human and everything they do
 - Motor planning
 - Language
 - Prediction
 - Learning


1950 Turky

1956 - AI coined

Dartmouth Conference.

Cognitive revolution

- Art
- Storytelling
- Humor
- Emotions



- **Early successes: logic, symbolic reasoning**

- Late 1950s, early 1960s
- Theorem provers
- Checkers and chess
- Why? (well-defined)

- **Knowledge-based AI: 1970s-1980s**

- Humans don't solve all problems from scratch
- Model humans and their processes
- Expert systems (e.g., medical diagnosis)
- Rule-based systems
- Brittle

- **AI Winter: 1987-1993**

- Rebirth with a dose of realism (?)

- Machine learning revolution


↳ Data
↳ Better Statistics
↳ Faster Machines

Weak AI
"Act Rationally"
Agents

Agents

Wednesday, August 24, 2016 8:47 AM

Anything that can perceive the environment through sensors and act on the env. through effectors



Agent function: A math. description of an agent's behaviors (response to the env.) mapping sensor percepts to actions through effectors.

$$F: S \rightarrow A$$


Agent Program: A concrete implementation of the agent function.

Table:

percept sequence	Action
---	⊖
---	⊖
---	⊖
---	⊖
---	⊖

Machine World

Vampires



2 locations
Robot
People

Actions:
Suck
Left
Right

Percepts	Action
(A, empty)	R
(A, people)	Suck
(B, empty)	L
(B, people)	S
(B, people)	?

- Table could be ∞ or very large
- Not enough HD
- Too big to author

Objective function: What we want the agent to achieve
Agent to evaluate all actions against objective
and choose the "best" action.

Rationality:

Choosing the "best" action given what you know.

What does an agent know?

percept history
built in knowledge
Actions

What does best mean? It depends

Vampires:

Maximize blood sucked.
 minimize movement
 Empty house
 Average emptiness over time

Mediodie job all the time
 flurry energy every so often.

Vampire Agent:

- sensors: perceive location, perceive people
- Actions: R, L, Suck
- Prior: map
- Objective fn: Greatest # of points over 1000 timesteps
- performance measure: 1 point for every empty room per timestep.

House starts empty


People enter house randomly

- occasionally check rooms - fast as possible

People enter, moving costs 1 point

You don't have a map

- explore
- exploit
- sensors unreliable
- effectors "
- Both



HW1 Due 9/25

Sun 11:55 PM

Max # points - 1 point/empty room/time step
- 1 point per move

- * Unreliable sensors
- * Unreliable effectors
- * sensors & effectors unreliable

Rational optimization \neq theoretical optimal

Expected Max \leftarrow

What type of agent we build depends on the type of environment.

Environments

Wednesday, August 24, 2016 8:44 AM

Observability

Fully Observable	Partially Observable
<ul style="list-style-type: none"> Can sense everything in the environment without error <p><i>checkers computer games logistics</i></p>	<ul style="list-style-type: none"> Can only see some info Two causes: <ul style="list-style-type: none"> Incomplete data ← Noise

Determinism

Deterministic	Stochastic <i>Randomness</i>
<ul style="list-style-type: none"> World only changes when agent acts World changes exactly as desired 	<ul style="list-style-type: none"> World can change at other times Randomness Uncertain effectors/sensors P.O. world is stochastic

Static

Static	Dynamic
<ul style="list-style-type: none">• World doesn't change when agent is deliberating	<ul style="list-style-type: none">• World changes faster than you can think

Discreteness

Discrete	Continuous
<ul style="list-style-type: none">• World broken up into a finite number of discrete chunks	<ul style="list-style-type: none">• Infinite number of chunks• Infinite gradations of values
<ul style="list-style-type: none">• Time• Actions• Percepts	


Episodic

Episodic	Sequential
<ul style="list-style-type: none">• History doesn't matter	<ul style="list-style-type: none">• History matters

Agents

Single agent	Multi agent
<ul style="list-style-type: none">• You are the only agent in the world	<ul style="list-style-type: none">• Other agents in the world that are autonomous from you<ul style="list-style-type: none">• Cooperative• Competitive

	Hanoi	Solitaire	Chess	Pool	Poker	Driving
Observability?	F.O.	P.O.	F.O.	F.O.	P.O.	P.O.
Determinism?	Det	Det /sto	Det/sto	sto	sto	sto.
Episodic?	Seq/Ep	Ep/seq	Seq.	Seq.	Seq.	Seq.
Static?	St.	St.	St.	St.	St.	Dyn.
Discrete?	Disc.	Disc	Disc.	Cont.	Disc.	Cont.
Agents?	S	S	M	M	M	M/S



Easiest? F.O., Det, stat, disc, ep, single-agent

Hardest? P.O., sto, Dyn, cont, seq, multi-agent

Real World

Part I: "Classical"

- Search & Problem Solving
- Planning
- Constraint Satisfaction
- Logical Agents ←

Part II: Stochastic

- Problem solving
- Learning